

# NLP 202: Structured Perceptron

---

Jeffrey Flanigan

Winter 2023


University of California Santa Cruz

[jmflanig@ucsc.edu](mailto:jmflanig@ucsc.edu)

- We have seen the simplest method for dependency parsing: *transition-based dependency parsing*
- Next simplest method: *graph-based dependency parsing with the structured perceptron*

I first need to introduce a new learning algorithm: the **Perceptron algorithm**

# The Perceptron



*Psychological Review*  
Vol. 65, No. 6, 1958

## THE PERCEPTRON: A PROBABILISTIC MODEL FOR INFORMATION STORAGE AND ORGANIZATION IN THE BRAIN<sup>1</sup>


F. ROSENBLATT

*Cornell Aeronautical Laboratory*

# The Perceptron algorithm

- Rosenblatt 1958
  - (Though there were some hints of a similar idea earlier, eg: Agmon 1954)
- The goal is to find a separating hyperplane
  - For separable data, guaranteed to find one
- An online algorithm
  - Processes one example at a time
- Several variants exist
  - We will see these briefly at towards the end

## Learning setup: Binary classification



- Training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in R^n, y_i \in \{-1, 1\}$
- Make predictions: 1 if  $\mathbf{w}^T \mathbf{x}_i \geq 0$ ,  $-1$  otherwise Another way to write this:  $\text{sign}(\mathbf{w}^T \mathbf{x})$
- Want to learn the weights  $\mathbf{w}$

Note: by include a constant feature of 1, the bias term can be “folded-in” to  $\mathbf{w}$

# Perceptron Algorithm

- **Decision rule:**  $\text{sign}(\mathbf{w}^T \mathbf{x}_i)$
- **Learning rule:** If incorrect:
  - if positive,  $\mathbf{w} \leftarrow \mathbf{w} + \mathbf{x}_i$   
if negative,  $\mathbf{w} \leftarrow \mathbf{w} - \mathbf{x}_i$   
Another way to write this:  $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i, y_i \in \{-1, 1\}$
- **Guaranteed to eventually correctly classify the data if the data are linearly separable**

# Perceptron Algorithm

- **Decision rule:**  $\text{sign}(\mathbf{w}^T \mathbf{x}_i)$
- **Learning rule:** If incorrect:
  - if positive,  $\mathbf{w} \leftarrow \mathbf{w} + \mathbf{x}_i$   
if negative,  $\mathbf{w} \leftarrow \mathbf{w} - \mathbf{x}_i$   
Another way to write this:  $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$ ,  $y_i \in \{-1, 1\}$
- **Guaranteed to eventually correctly classify the data if the data are linearly separable**

1. Initialize  $\mathbf{w} = \mathbf{0} \in R^n$
2. For epoch in  $1 \dots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \neq \text{sign}(\mathbf{w}^T \mathbf{x})$  (shorthand:  $y_i \mathbf{w}^T \mathbf{x}_i \leq 0$ ), then:  
Update  $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$
3. Return  $\mathbf{w}$

# Perceptron Algorithm

- **Decision rule:**  $\text{sign}(\mathbf{w}^T \mathbf{x}_i)$
  - **Learning rule:** If incorrect:
    - if positive,  $\mathbf{w} \leftarrow \mathbf{w} + \mathbf{x}_i$   
if negative,  $\mathbf{w} \leftarrow \mathbf{w} - \mathbf{x}_i$   
Another way to write this:  $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$ ,  $y_i \in \{-1, 1\}$
  - **Guaranteed to eventually correctly classify the data if the data are linearly separable**
1. Initialize  $\mathbf{w} = \mathbf{0} \in R^n$
  2. For epoch in  $1 \dots T$ :
    1. Shuffle the data
    2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
      - If  $y_i \neq \text{sign}(\mathbf{w}^T \mathbf{x})$  (shorthand:  $y_i \mathbf{w}^T \mathbf{x}_i \leq 0$ ), then:  
Update  $\mathbf{w} \leftarrow \mathbf{w} + r y_i \mathbf{x}_i$  Can also have a stepsize (learning rate)  $r$ . If constant, can set to 1
  3. Return  $\mathbf{w}$



# Intuition behind the update

Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r\mathbf{x}_i$

Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r\mathbf{x}_i$

Suppose we have made a mistake on a positive example

That is,  $y = +1$  and  $\mathbf{w}_t^T \mathbf{x} \leq 0$

Call the **new weight vector**  $\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{x}$  (say  $r = 1$ )


The **new dot product** is  $\mathbf{w}_{t+1}^T \mathbf{x} = (\mathbf{w}_t + \mathbf{x})^T \mathbf{x} = \mathbf{w}_t^T \mathbf{x} + \mathbf{x}^T \mathbf{x} \geq \mathbf{w}_t^T \mathbf{x}$

*For a positive example, the Perceptron update will increase the score assigned to the same input*

Similar reasoning for negative examples


Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r\mathbf{x}_i$   
Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r\mathbf{x}_i$

# Geometry of the perceptron update




Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r\mathbf{x}_i$   
Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r\mathbf{x}_i$

# Geometry of the perceptron update



Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r\mathbf{x}_i$   
Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r\mathbf{x}_i$


# Geometry of the perceptron update



For a mistake on a **positive** example

Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r\mathbf{x}_i$   
Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r\mathbf{x}_i$


# Geometry of the perceptron update



For a mistake on a **positive** example

Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r\mathbf{x}_i$   
Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r\mathbf{x}_i$


# Geometry of the perceptron update



For a mistake on a **positive** example

Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r\mathbf{x}_i$   
Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r\mathbf{x}_i$


# Geometry of the perceptron update



For a mistake on a **positive** example

Mistake on positive:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + r\mathbf{x}_i$   
Mistake on negative:  $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - r\mathbf{x}_i$


# Geometry of the perceptron update




For a mistake on a **positive** example



# Geometry of the perceptron update




# Geometry of the perceptron update




For a mistake on a **negative** example

# Geometry of the perceptron update




For a mistake on a **negative** example

# Geometry of the perceptron update




For a mistake on a **negative** example

# Geometry of the perceptron update



For a mistake on a **negative** example

# Geometry of the perceptron update



For a mistake on a **negative**  
example

Variants of the Perceptron: voting and averaging

# 1. The “standard” algorithm

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathbb{R}^n, y_i \in \{-1, 1\}$

1. Initialize  $\mathbf{w} = \mathbf{0} \in \mathbb{R}^n$
2. For epoch in  $1 \cdots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \mathbf{w}^T \mathbf{x}_i \leq 0$ , then:
      - update  $\mathbf{w} \leftarrow \mathbf{w} + r y_i \mathbf{x}_i$
3. Return  $\mathbf{w}$

**Prediction** on a new example with features  $\mathbf{x}$ :  $\text{sgn}(\mathbf{w}^T \mathbf{x})$



# 1. The “standard” algorithm

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathbb{R}^n, y_i \in \{-1, 1\}$

1. Initialize  $\mathbf{w} = \mathbf{0} \in \mathbb{R}^n$
2. For epoch in  $1 \dots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \mathbf{w}^T \mathbf{x}_i \leq 0$ , then:
      - update  $\mathbf{w} \leftarrow \mathbf{w} + r y_i \mathbf{x}_i$
3. Return  $\mathbf{w}$

$T$  is a **hyper-parameter** to the algorithm

Another way of writing that there is an error

**Prediction** on a new example with features  $\mathbf{x}$ :  $\text{sgn}(\mathbf{w}^T \mathbf{x})$

## 2. Voting and Averaging

- So far: We return the final weight vector
- Voted perceptron
  - Remember every weight vector in your sequence of updates.
  - At final prediction time, each weight vector gets to vote on the label. The number of votes it gets is the number of iterations it survived before being updated
  - Comes with strong theoretical guarantees about generalization, impractical because of storage issues

## 2. Voting and Averaging

- So far: We return the final weight vector
- Voted perceptron
  - Remember every weight vector in your sequence of updates.
  - At final prediction time, each weight vector gets to vote on the label. The number of votes it gets is the number of iterations it survived before being updated
  - Comes with strong theoretical guarantees about generalization, impractical because of storage issues
- Averaged perceptron
  - Instead of using all weight vectors, use the average weight vector (i.e longer surviving weight vectors get more say)
  - More practical alternative and widely used

# Averaged Perceptron

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathbb{R}^n, y_i \in \{-1, 1\}$

1. Initialize  $\mathbf{w} = \mathbf{0} \in \mathbb{R}^n$  and  $\mathbf{a} = \mathbf{0} \in \mathbb{R}^n$
2. For epoch in  $1 \cdots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \mathbf{w}^T \mathbf{x}_i \leq 0$ , then:
      - update  $\mathbf{w} \leftarrow \mathbf{w} + r y_i \mathbf{x}_i$
    - $\mathbf{a} \leftarrow \mathbf{a} + \mathbf{w}$
3. Return  $\mathbf{a}$

**Prediction** on a new example with features  $\mathbf{x}$ :  $\text{sgn}(\mathbf{a}^T \mathbf{x})$

# Averaged Perceptron

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathbb{R}^n, y_i \in \{-1, 1\}$

1. Initialize  $\mathbf{w} = \mathbf{0} \in \mathbb{R}^n$  and  $\mathbf{a} = \mathbf{0} \in \mathbb{R}^n$
2. For epoch in  $1 \cdots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \mathbf{w}^T \mathbf{x}_i \leq 0$ , then:
      - update  $\mathbf{w} \leftarrow \mathbf{w} + r y_i \mathbf{x}_i$
    - $\mathbf{a} \leftarrow \mathbf{a} + \mathbf{w}$
3. Return  $\mathbf{a}$

This is the simplest version of the averaged perceptron

There are some easy programming tricks to make sure that  $\mathbf{a}$  is also updated even when there is an error

**Prediction** on a new example with features  $\mathbf{x}$ :  $\text{sgn}(\mathbf{a}^T \mathbf{x})$

# Averaged Perceptron

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathbb{R}^n, y_i \in \{-1, 1\}$

1. Initialize  $\mathbf{w} = \mathbf{0} \in \mathbb{R}^n$  and  $\mathbf{a} = \mathbf{0} \in \mathbb{R}^n$
2. For epoch in  $1 \cdots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \mathbf{w}^T \mathbf{x}_i \leq 0$ , then:
      - update  $\mathbf{w} \leftarrow \mathbf{w} + r y_i \mathbf{x}_i$
    - $\mathbf{a} \leftarrow \mathbf{a} + \mathbf{w}$
3. Return  $\mathbf{a}$

This is the simplest version of the averaged perceptron

There are some easy programming tricks to make sure that  $\mathbf{a}$  is also updated only when there is an error

*If you want to use the Perceptron algorithm, use averaging*

**Prediction** on a new example with features  $\mathbf{x}$ :  $\text{sgn}(\mathbf{a}^T \mathbf{x})$

# Analysis: Perceptron

## Perceptron Mistake Bound

**Theorem 0.1** (Block (1962), Novikoff (1962)).


Given dataset:  $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$ .

Suppose:

1. Finite size inputs:  $\|\mathbf{x}^{(i)}\| \leq R$
2. Linearly separable data:  $\exists \boldsymbol{\theta}^*$  s.t.  $\|\boldsymbol{\theta}^*\| = 1$  and  $y^{(i)}(\boldsymbol{\theta}^* \cdot \mathbf{x}^{(i)}) \geq \gamma, \forall i$

Then: The number of mistakes made by the Perceptron algorithm on this dataset is

$$k \leq (R/\gamma)^2$$



# Standard Perceptron Algorithm (Binary Case)

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in R^n, y_i \in \{-1, 1\}$

1. Initialize  $\mathbf{w} = \mathbf{0} \in R^n$
2. For epoch in  $1 \dots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \mathbf{w}^T \mathbf{x}_i \leq 0$ , then:  
Update  $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$
3. Return  $\mathbf{w}$

Prediction on a new example with features  $\mathbf{x}$ :  $\text{sign}(\mathbf{w}^T \mathbf{x})$



What about multiclass?

# Multiclass Perceptron

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathcal{X}$ ,  $y_i \in L$  and linear scoring function  $\mathbf{w}^T \mathbf{f}(\mathbf{x}_i, y)$

1. Initialize  $\mathbf{w} = \mathbf{0} \in R^n$
2. For epoch in  $1 \dots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \neq \hat{y} = \operatorname{argmax}_y \mathbf{w}^T \mathbf{f}(\mathbf{x}_i, y)$ , then:  
Update  $\mathbf{w} \leftarrow \mathbf{w} + \mathbf{f}(\mathbf{x}_i, y_i) - \mathbf{f}(\mathbf{x}_i, \hat{y})$
3. Return  $\mathbf{w}$

Prediction on a new example  $\mathbf{x}$ :  $\hat{y} = \operatorname{argmax}_y \mathbf{w}^T \mathbf{f}(\mathbf{x}_i, y)$

# Multiclass Perceptron

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathcal{X}$ ,  $y_i \in L$  and linear scoring function  $\mathbf{w}^T \mathbf{f}(\mathbf{x}_i, y)$

1. Initialize  $\mathbf{w} = \mathbf{0} \in R^n$
2. For epoch in  $1 \dots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - If  $y_i \neq \hat{y} = \operatorname{argmax}_y \mathbf{w}^T \mathbf{f}(\mathbf{x}_i, y)$ , then *(this check is not necessary, because if same, then update cancels)*:  
Update  $\mathbf{w} \leftarrow \mathbf{w} + \mathbf{f}(\mathbf{x}_i, y_i) - \mathbf{f}(\mathbf{x}_i, \hat{y})$
3. Return  $\mathbf{w}$

Prediction on a new example  $\mathbf{x}$ :  $\hat{y} = \operatorname{argmax}_y \mathbf{w}^T \mathbf{f}(\mathbf{x}, y)$

# Multiclass Perceptron

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathcal{X}$ ,  $y_i \in L$  and linear scoring function  $\mathbf{w}^T \mathbf{f}(\mathbf{x}_i, y)$

1. Initialize  $\mathbf{w} = \mathbf{0} \in R^n$
2. For epoch in  $1 \dots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - $\hat{y} = \operatorname{argmax}_y \mathbf{w}^T \mathbf{f}(\mathbf{x}_i, y)$
    - Update  $\mathbf{w} \leftarrow \mathbf{w} + \mathbf{f}(\mathbf{x}_i, y_i) - \mathbf{f}(\mathbf{x}_i, \hat{y})$
3. Return  $\mathbf{w}$

Prediction on a new example  $\mathbf{x}$ :  $\hat{y} = \operatorname{argmax}_y \mathbf{w}^T \mathbf{f}(\mathbf{x}_i, y)$

Can we use the perceptron for **structured outputs**?

For example, sequence labeling and dependency parsing?

# Structured Perceptron (with linear scoring function)

Given a training set  $D = \{(\mathbf{x}_i, y_i)\}$  where  $\mathbf{x}_i \in \mathcal{X}$ ,  $y_i \in \mathcal{Y}(\mathbf{x}_i)$  (*output space depends on input*) and linear scoring function  $\mathbf{w}^T \mathbf{f}(\mathbf{x}_i, \mathbf{y})$

1. Initialize  $\mathbf{w} = \mathbf{0} \in R^n$
2. For epoch in  $1 \dots T$ :
  1. Shuffle the data
  2. For each training example  $(\mathbf{x}_i, y_i) \in D$ :
    - Make prediction  $\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y}} \mathbf{w}^T \mathbf{f}(\mathbf{x}_i, \mathbf{y})$
    - Update  $\mathbf{w} \leftarrow \mathbf{w} + \mathbf{f}(\mathbf{x}_i, y_i) - \mathbf{f}(\mathbf{x}_i, \hat{\mathbf{y}})$
3. Return  $\mathbf{w}$

Prediction on a new example  $\mathbf{x}$ :  $\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y}} \mathbf{w}^T \mathbf{f}(\mathbf{x}_i, \mathbf{y})$

## Another view of the Perceptron

It turns out, the Perceptron algorithm is exactly stochastic (sub)-gradient descent on a loss function!

The Perceptron minimizes:

$$\min_{\mathbf{w}} \sum_{i=1}^N \left( \max_{y' \in \mathcal{Y}} \mathbf{w} \cdot f(\mathbf{x}_i, y') \right) - \mathbf{w} \cdot f(\mathbf{x}_i, y_i)$$

# Minimizing the Perceptron Loss: The Perceptron Algorithm

Perceptron solves the following minimization problem:

$$\min_{\mathbf{w}} \sum_{i=1}^N \left( \max_{y' \in \mathcal{Y}} \mathbf{w} \cdot f(\mathbf{x}_i, y') \right) - \mathbf{w} \cdot f(\mathbf{x}_i, y_i)$$

Stochastic subgradient descent (SSGD) with stepsize  $\alpha$  on the above is the **Perceptron** algorithm!

- For  $i \in \{1, \dots, N\}$ :
  - Shuffle the training data, and for each example  $i$  do the following update:
    - $\hat{y}_i \leftarrow \operatorname{argmax}_{y \in \mathcal{Y}} \mathbf{w} \cdot f(\mathbf{x}_i, y)$
    - $\mathbf{w} \leftarrow \mathbf{w} - \alpha (f(\mathbf{x}_i, \hat{y}) - f(\mathbf{x}_i, y_i))$



## Minimizing Perceptron Loss

$$\left( \max_{y' \in \mathcal{Y}} \mathbf{w} \cdot f(\mathbf{x}, y') \right) - \mathbf{w} \cdot f(\mathbf{x}, y)$$


## Minimizing Perceptron Loss

$$\left( \max_{y' \in \mathcal{Y}} \mathbf{w} \cdot f(\mathbf{x}, y') \right) - \mathbf{w} \cdot f(\mathbf{x}, y)$$

When two labels are *tied*, the function is not differentiable.

Solution: (stochastic) subgradient descent!

# Subgradient



$c$  is a **subgradient** of  $f(x)$  at  $x_0$  if  $\forall x$ :

$$f(x) - f(x_0) \geq c(x - x_0)$$

(Only defined for convex functions.)

# Gradient of Perceptron Loss

For the  $i$ th data point the subgradient is:

$$\begin{aligned} g_i^j &= \frac{\partial}{\partial \mathbf{w}_j} \left[ \left( \max_{y' \in \mathcal{Y}} \mathbf{w} \cdot f(x_i, y') \right) - \mathbf{w} \cdot f(x_i, y_i) \right] \\ &= f^j(x_i, \hat{y}) - f^j(x_i, y_i) \end{aligned}$$

where

$$\hat{y} = \operatorname{argmax}_{y' \in \mathcal{Y}} \mathbf{w} \cdot f(x_i, y')$$

# Minimizing the Perceptron Loss: The Perceptron Algorithm

Perceptron solves the following minimization problem:

$$\min_{\mathbf{w}} \sum_{i=1}^N \left( \max_{y' \in \mathcal{Y}} \mathbf{w} \cdot f(\mathbf{x}_i, y') \right) - \mathbf{w} \cdot f(\mathbf{x}_i, y_i)$$

Stochastic subgradient descent (SSGD) with stepsize  $\alpha$  on the above is the **Perceptron** algorithm!

- For  $i \in \{1, \dots, N\}$ :
  - Shuffle the training data, and for each example  $i$  do the following update:
    - $\hat{y}_i \leftarrow \operatorname{argmax}_{y \in \mathcal{Y}} \mathbf{w} \cdot f(\mathbf{x}_i, y)$
    - $\mathbf{w} \leftarrow \mathbf{w} - \alpha (f(\mathbf{x}_i, \hat{y}) - f(\mathbf{x}_i, y_i))$


Proof the the Perceptron mistake bound (not covered in class).

# Analysis: Perceptron

## Perceptron Mistake Bound

**Guarantee:** If data has margin  $\gamma$  and all points inside a ball of radius  $R$ , then Perceptron makes  $\leq (R/\gamma)^2$  mistakes.

(Normalized margin: multiplying all points by 100, or dividing all points by 100, doesn't change the number of mistakes; algo is invariant to scaling.)



# Analysis: Perceptron

## Perceptron Mistake Bound

**Theorem 0.1** (Block (1962), Novikoff (1962)).


Given dataset:  $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$ .

Suppose:

1. Finite size inputs:  $\|\mathbf{x}^{(i)}\| \leq R$
2. Linearly separable data:  $\exists \boldsymbol{\theta}^*$  s.t.  $\|\boldsymbol{\theta}^*\| = 1$  and  $y^{(i)}(\boldsymbol{\theta}^* \cdot \mathbf{x}^{(i)}) \geq \gamma, \forall i$

Then: The number of mistakes made by the Perceptron algorithm on this dataset is

$$k \leq (R/\gamma)^2$$






# Analysis: Perceptron

## Proof of Perceptron Mistake Bound:

We will show that there exist constants A and B s.t.

$$Ak \leq ||\boldsymbol{\theta}^{(k+1)}|| \leq B\sqrt{k}$$



# Analysis: Perceptron

## Proof of Perceptron Mistake Bound:

Part 1: for some  $A$ ,  $Ak \leq \|\theta^{(k+1)}\|$

$$\theta^{(k+1)} \cdot \theta^* = (\theta^{(k)} + y^{(i)} \mathbf{x}^{(i)}) \theta^*$$

by Perceptron algorithm update

$$= \theta^{(k)} \cdot \theta^* + y^{(i)} (\theta^* \cdot \mathbf{x}^{(i)})$$

$$\geq \theta^{(k)} \cdot \theta^* + \gamma$$

by assumption

$$\Rightarrow \theta^{(k+1)} \cdot \theta^* \geq k\gamma$$

by induction on  $k$  since  $\theta^{(1)} = \mathbf{0}$

$$\Rightarrow \|\theta^{(k+1)}\| \geq k\gamma$$

since  $\|\mathbf{w}\| \times \|\mathbf{u}\| \geq \mathbf{w} \cdot \mathbf{u}$  and  $\|\theta^*\| = 1$

Cauchy-Schwartz inequality

# Analysis: Perceptron

## Proof of Perceptron Mistake Bound:

Part 2: for some  $B$ ,  $\|\boldsymbol{\theta}^{(k+1)}\| \leq B\sqrt{k}$

$$\|\boldsymbol{\theta}^{(k+1)}\|^2 = \|\boldsymbol{\theta}^{(k)} + y^{(i)}\mathbf{x}^{(i)}\|^2$$

by Perceptron algorithm update

$$= \|\boldsymbol{\theta}^{(k)}\|^2 + (y^{(i)})^2 \|\mathbf{x}^{(i)}\|^2 + 2y^{(i)}(\boldsymbol{\theta}^{(k)} \cdot \mathbf{x}^{(i)})$$

$$\leq \|\boldsymbol{\theta}^{(k)}\|^2 + (y^{(i)})^2 \|\mathbf{x}^{(i)}\|^2$$

since  $k$ th mistake  $\Rightarrow y^{(i)}(\boldsymbol{\theta}^{(k)} \cdot \mathbf{x}^{(i)}) \leq 0$

$$= \|\boldsymbol{\theta}^{(k)}\|^2 + R^2$$

since  $(y^{(i)})^2 \|\mathbf{x}^{(i)}\|^2 = \|\mathbf{x}^{(i)}\|^2 = R^2$  by assumption and  $(y^{(i)})^2 = 1$

$$\Rightarrow \|\boldsymbol{\theta}^{(k+1)}\|^2 \leq kR^2$$

by induction on  $k$  since  $(\boldsymbol{\theta}^{(1)})^2 = 0$

$$\Rightarrow \|\boldsymbol{\theta}^{(k+1)}\| \leq \sqrt{k}R$$

# Analysis: Perceptron

## Proof of Perceptron Mistake Bound:

Part 3: Combining the bounds finishes the proof.

$$k\gamma \leq ||\boldsymbol{\theta}^{(k+1)}|| \leq \sqrt{k}R$$

$$\Rightarrow k \leq (R/\gamma)^2$$




The total number of mistakes  
must be less than this

## Log Loss and Perceptron Loss for $(\mathbf{x}, y)$

$$\text{log loss: } \left( \log \sum_{y' \in \mathcal{Y}} \exp \mathbf{w} \cdot f(\mathbf{x}, y') \right) - \mathbf{w} \cdot f(\mathbf{x}, y)$$

$$\text{Perceptron loss: } \left( \max_{y' \in \mathcal{Y}} \mathbf{w} \cdot f(\mathbf{x}, y') \right) - \mathbf{w} \cdot f(\mathbf{x}, y)$$



In **purple** is the hinge loss, in **blue** is the log loss; in **red** is the “zero-one” loss (error).